

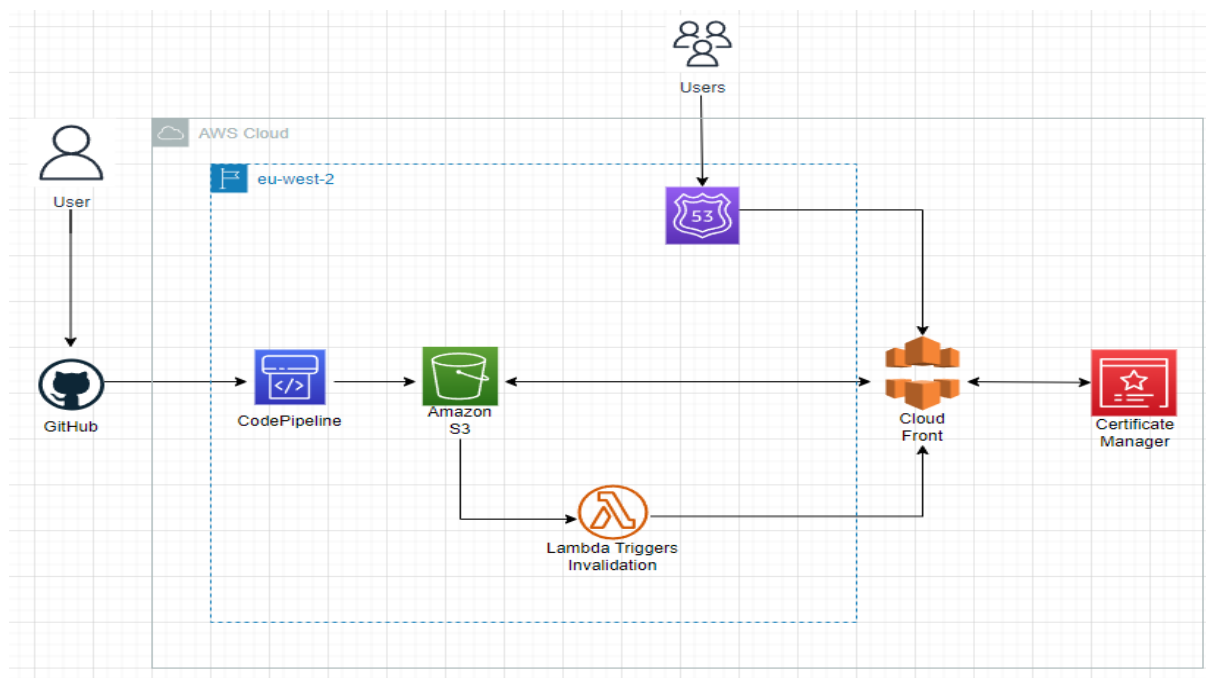
Static Website hosted on S3 with CI/CD (step-by-step guide)

Introduction:

This tutorial will guide you through the process of hosting a static website on Amazon S3 and making it accessible worldwide using Amazon CloudFront. You'll also learn how to secure your website using OAC to keep the bucket private, AWS Certificate Manager to protect and secure your website, and manage your domain with Amazon Route 53. Last but not least, we will use GitHub and AWS CodePipeline to automate every change we want to make to our website.

Static websites are a cost-effective and straightforward way to establish an online presence, whether it's for personal use, a blog, or a business. By following the steps in this guide, you can quickly host and manage your own static website with AWS's robust cloud services. This setup ensures speed, security, reliability, along with a custom domain and HTTPS support.

Here is an architecture on what you will be able to build by the end of this tutorial:



Prerequisites:

AWS Account: To follow the examples and implement the project, you will need an AWS account.

AWS Services: While detailed knowledge of AWS services is not required, having a general awareness of services like Amazon S3, Amazon CloudFront, AWS CodePipeline, Lambda, AWS Identity and Access Management (IAM), AWS Certificate Manager, and Amazon Route 53 will be helpful. This tutorial will guide you through the setup and usage of these services.

GitHub Account: You should have a GitHub account for version control and collaborative development. Basic knowledge of Git for version control is a plus but not strictly required, as this tutorial may provide some Git basics.

Visual Studio Code: Familiarity with Visual Studio Code (VS Code) is recommended, including how to open, edit, and save files using this code editor.

Web Development Basics (Optional): Some experience with web development technologies (HTML, CSS, JavaScript) is helpful but not mandatory, as this tutorial will primarily focus on the AWS and DevOps aspects.

Audience:

This tutorial is designed to accommodate a broad range of users, from beginners to advanced AWS practitioners. Whether you're new to AWS services or an experienced cloud engineer, you'll find valuable information and guidance tailored to your level of expertise.

Beginners: If you're just starting your AWS journey, this tutorial will provide step-by-step instructions and explanations to help you grasp the fundamentals and build a strong foundation in cloud static hosting.

Intermediate Users: If you have some experience with AWS but want to expand your knowledge and enhance your skills, you'll find in-depth insights, best practices, and tips for optimizing your projects.

Advanced Users: For experienced AWS professionals, this tutorial offers advanced configurations, automation strategies, and techniques to streamline your workflow and achieve a higher level of efficiency in deploying static websites.

No matter where you fall on the spectrum of AWS expertise, this tutorial is designed to empower you with the knowledge and tools necessary to successfully host static websites and implement CI/CD pipelines in AWS.

PS: You can find all the files you need for this tutorial on <https://github.com/Paul-Ciurean/Tutorial-S3-with-CICD.git>

Step 1: Create our buckets

1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.

2. Choose **Create bucket**.

3. Enter the **Bucket name** for example, **example.com**.(main bucket)

4. Choose the Region where you want to create the bucket.

Choose a Region that is geographically close to you to minimize latency and costs, or to address regulatory requirements. The Region that you choose determines your Amazon S3 website endpoint. For more information, see [Website endpoints](#).

5. To accept the default settings and create the bucket, choose **Create**.

6. Follow the above steps to create a 2nd bucket with the name **www.example.com**

<input type="radio"/>	www.paulciurean.com	Europe (London) eu-west-2	<u>Bucket and objects not public</u>
<input type="radio"/>	paulciurean.com	Europe (London) eu-west-2	<u>Bucket and objects not public</u>

We will keep our buckets private, as we will use Origin Access Control(OAC) to restrict users to access our buckets only through Cloud Front distribution.

To make sure that the users will be redirected to the main bucket, we will have to enable static website on our 2nd bucket.

1. We do that by going into our 2nd bucket, we chose **Properties**, and scroll all the way down to the bottom of the page.

2. We click **Edit** and **Enable** for **Static website hosting**.

3. On the section **Hosting type** we chose **Redirect request for an object** and on the **Host name** we specify our main bucket.

4. On **Protocol** we select **HTTPS**.

5. **Save changes**.

Your page should look like this:

Edit static website hosting [Info](#)

Static website hosting

Use this bucket to host a website or redirect requests. [Learn more](#)

Static website hosting

☐ Disable

☒ Enable

Hosting type

☐ Host a static website
Use the bucket endpoint as the web address. [Learn more](#)

☒ Redirect requests for an object
Redirect requests to another bucket or domain. [Learn more](#)

Host name

Target bucket website address or personal domain

Protocol - *Optional*

☐ none

☐ http

☒ https

[Cancel](#) [Save changes](#)

Uploading Website files:

You should upload your Website files in the main bucket(paulciurean.com). For the Static Website you should have an **index.html** file and **error.html**(optional) file. In my **index.html**

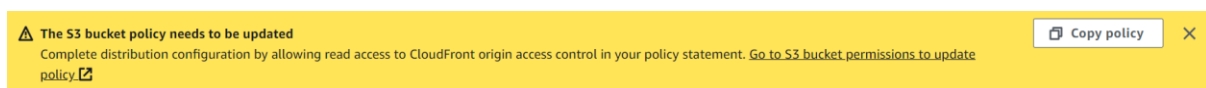
Files and folders (2 Total, 497.0 B)						
<input type="text" value="Find by name"/>						
Name	Folder	Type	Size	Status		
error.html	-	text/html	208.0 B	✓ Succeeded		
index.html	-	text/html	289.0 B	✓ Succeeded		

In my **index.html** I've got a basic webpage created in HTML.

Step 2: Set up Cloud Front Distribution

1. In **Cloud Front** on **Distribution** page we **Create Distribution**.
2. On the **Origin** page for **Origin domain** we select our main S3 bucket.
3. We give a **Name** to our origin.
4. **Origin Access** we chose **Origin access control settings**.
5. **Origin access control** we create a control setting.
 - 5.1. In the new window we give a **Name**.
 - 5.2. **Origin type** we select **S3**.
 - 5.3. We keep the default option for the rest.
6. On the page **Default cache behaviour** at the section **Viewer protocol policy** we select **Redirect HTTP to HTTPS**. We keep the default option for the rest.
7. **Web Application Firewall**: We **Do not enable security protection**.
8. **Settings** page:
 - 8.1 **Alternate domain name** -> **Add item**: add the custom domain names. (**paulciurean.com** and **www.paulciurean.com**)
 - 8.2 **Custom SSL certificate**: we request a certificate.
 - 8.3 **Default root object**: **index.html** (name of our index file).
9. **Create distribution**.

When popped with the following message:



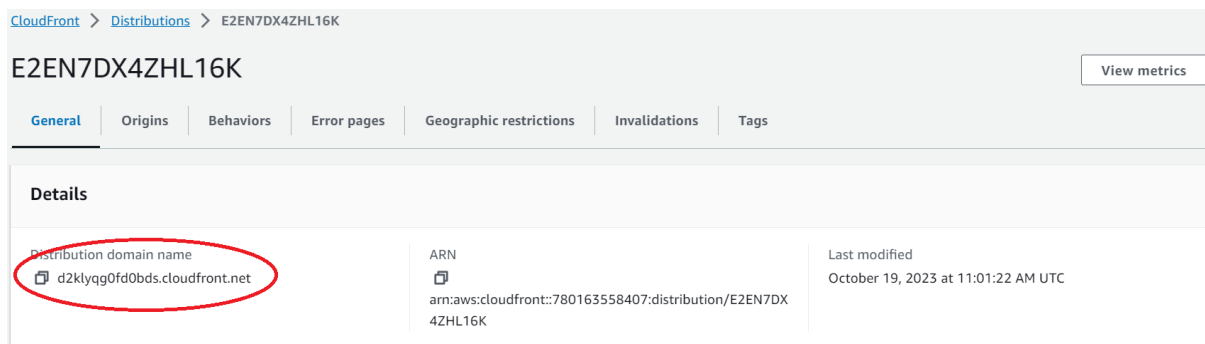
Copy the policy, we will have to complete our distribution configuration by allowing read access to CloudFront OAC in our S3 bucket policy statement.

To do that follow the next steps:

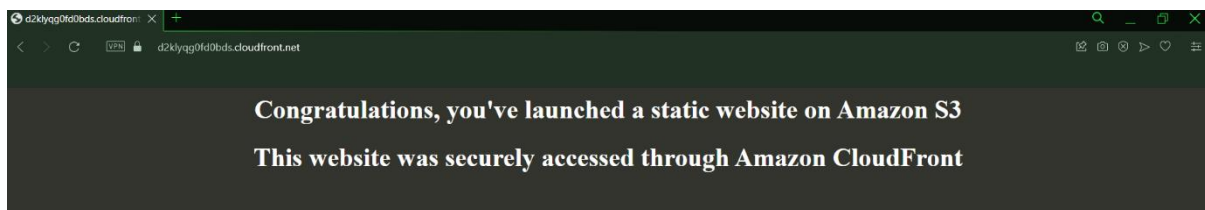
1. Go to our main bucket.
2. In **Permissions** scroll down to **Bucket policy** and click **Edit**.
3. Paste the policy you copied from the CloudFront distribution.

PS: If you did not copy the policy after creating the distribution, you can find it in CloudFront -> Distribution -> Origins -> select Origins name and click Edit. On the following page you can find the policy.

To test if our website is working, you can copy the **Distribution domain name** from your distribution:



And paste it in a new browser. My page looks like this:



PS: If it doesn't work that means the CloudFront distribution is not deployed yet. Give it few minutes and refresh the page.

Step 3: Domain and DNS Configuration

Register a Domain (if you don't already have one).

Set up DNS Records:

1. In **Route 53** on the **Hosted zone** page, **Create hosted zone**.
2. **Select the name of your website (paulciurean.com).**
3. **Type: Public hosted zone.**
4. **Create hosted zone.**

By default there will be 2 record names.

We now create 2 other records: 1 for paulciurean.com and 1 for www.paulciurean.com

5. Click **Create record**.
6. Click on **Switch to quick create**.
7. **Record name:** leave it blank. (when you create 2nd record you will have to type **www**).
8. For **Record type** chose **A**.
9. Above **Value** click **Alias**.
10. **Route traffic to:** Alias to CloudFront distribution.
11. Chose your distribution created in **Step 2**.
12. Click **Create records**.
13. Follow the previous steps to create a record for "www".

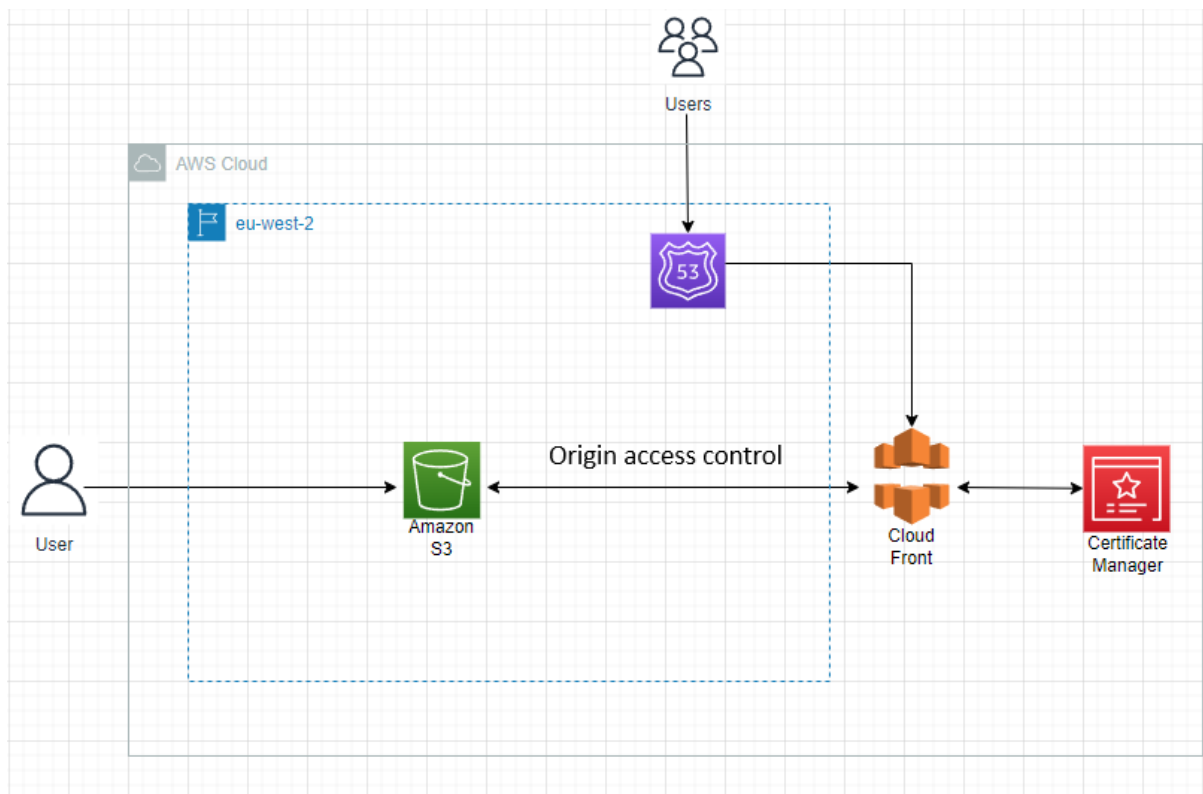
You are now ready to test the website by opening a browser and browse the following URLs:

http://your-domain-name, for example, example.com – Displays the index document in the your-domain-name bucket

http://www.your-domain-name for example, www.example.com – Redirects your request to the your-domain-name bucket

PS: If it doesn't work that's because usually DNS changes will propagate within a few hours, but it can take up to 48 hours for everything to propagate across the Internet.

So far we achieved the following :



Up to this point, this is in detail tutorial on how to set up a static website using S3, keep the bucket private and secure using OAC, make it accessible worldwide using CloudFront for low latency, secure and protect your website with Certificate Manager, and manage your domain with Amazon Route 53.

If we want to make any changes to our website, we have to change the files (**index.html** and others such as **style.css** and **script.js** in case we want to add style to our website, not just keep it a blank page), delete and upload the files again in S3 manually, go to CloudFront and create an invalidation to remove the cache so the users can see the up-to-date site.

Up next we will continue building a CI/CD pipeline to automate the process of making any changes to our website.

Step 4: Using Visual Studio Code and Github

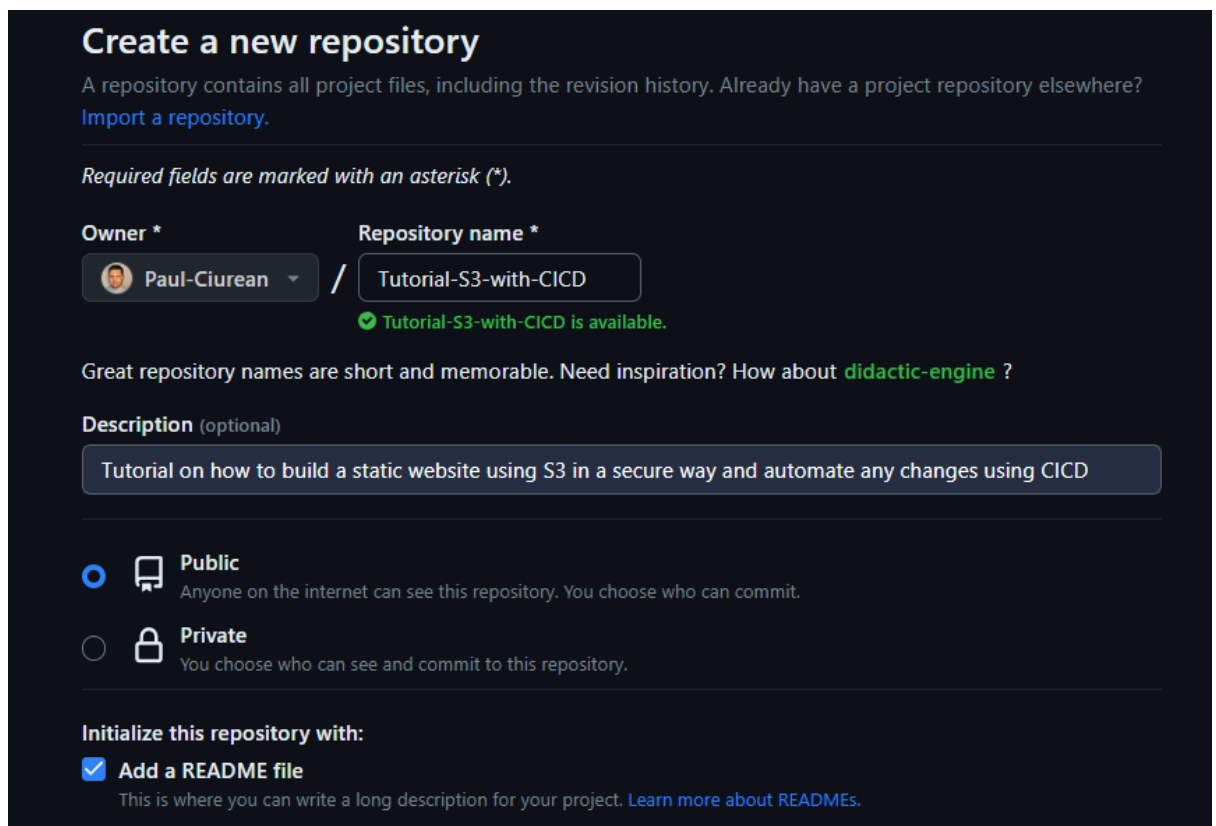
You should have a basic understanding on how to use VSC and github for the following part.

Here is a list of videos I watched to set up and get used to VSC and github:

https://www.youtube.com/watch?v=i_23KUAEtUM

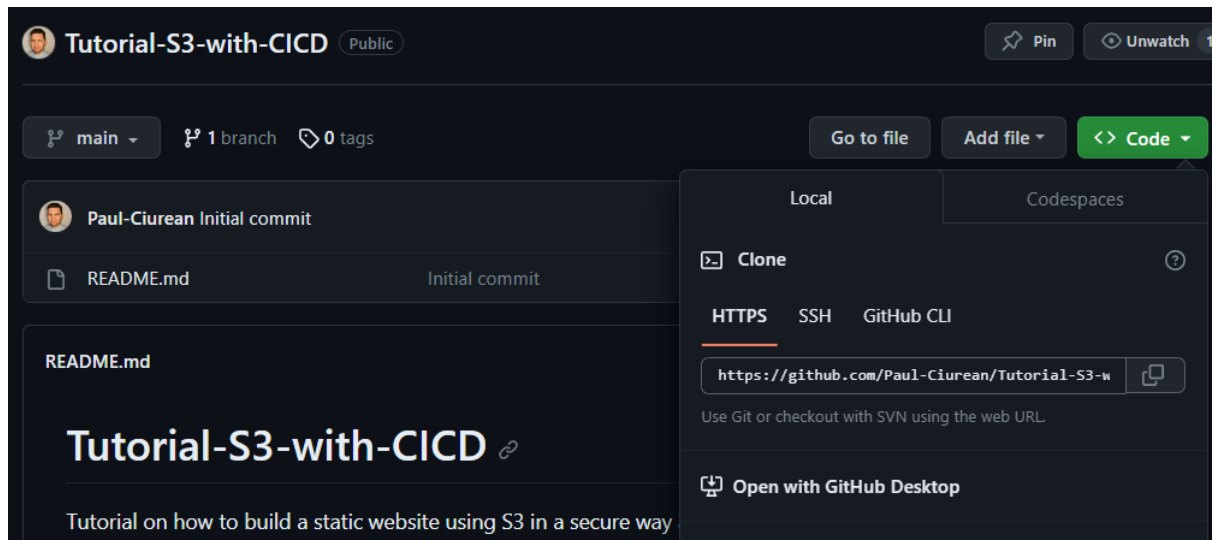
<https://www.youtube.com/watch?v=RG0j5yH7evk&t=1475s>

1. Open your GitHub account and create a new repository.

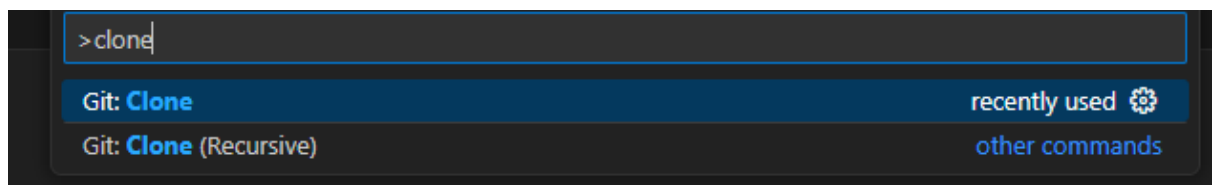


The screenshot shows the GitHub 'Create a new repository' interface. At the top, it says 'Create a new repository' and provides a brief explanation of what a repository is. Below this, there's a section for 'Required fields are marked with an asterisk (*)'. The 'Owner' field is set to 'Paul-Ciurean' and the 'Repository name' field is 'Tutorial-S3-with-CICD'. A green checkmark indicates that the name is available. The 'Description' field is optional and contains the text 'Tutorial on how to build a static website using S3 in a secure way and automate any changes using CICD'. Under the 'Visibility' section, 'Public' is selected, with a description: 'Anyone on the internet can see this repository. You choose who can commit.' The 'Private' option is also visible. At the bottom, under 'Initialize this repository with:', the 'Add a README file' checkbox is checked, with a note: 'This is where you can write a long description for your project. Learn more about READMEs.'


2. Once the repository is created open VSC and press **Shift+CTRL+P**
3. Copy the link from **Local -> HTTPS** in your clipboard.



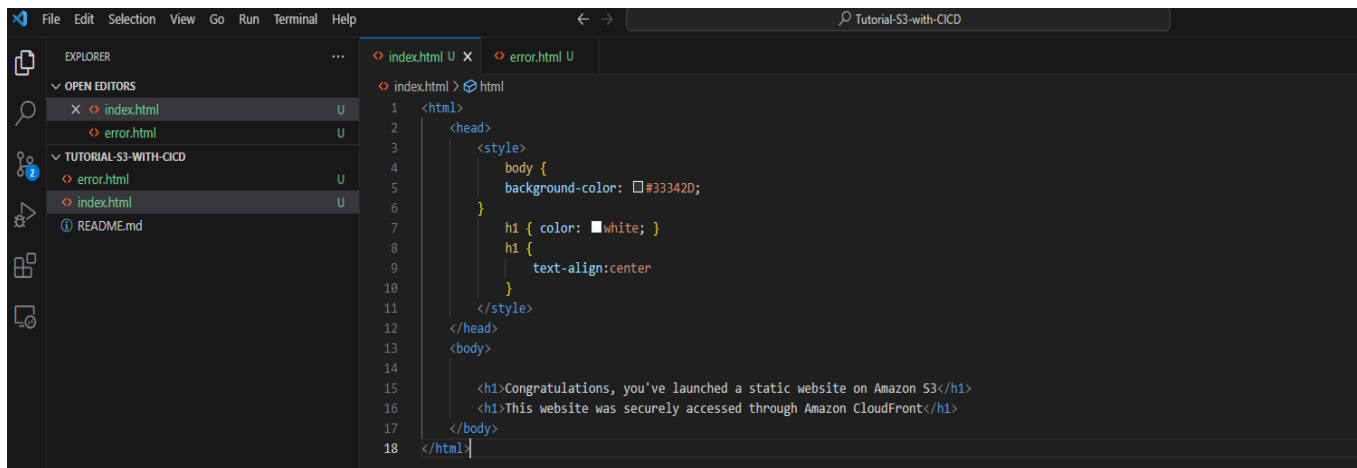
4. Type **clone** to clone the repository



5. Paste the link from Github website in the section above.
That will take you to a window where you will have to select a folder where to keep your files locally on your machine.
PS: This will only work after you connect the github account with your local machine. Follow the tutorials from the youtube videos from above.
6. You should now be in VSC with the repository open looking like this:




Here is the place where you can create your **index.html**, **styles.css**, and all the rest of the files for your website. In my repository you will be able to find the files I am using for this tutorial.




To commit the changes to the **github** repo into your online **github** account follow the next steps:

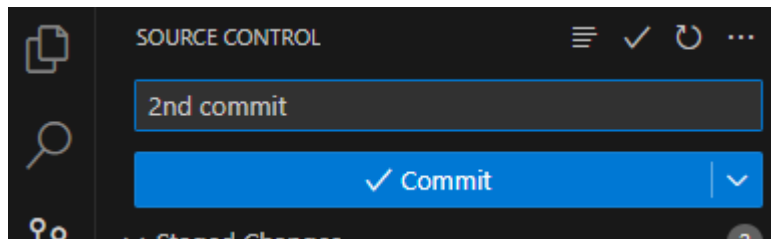
1. Click on source control on the left side of the page:



2. Stage all changes. Press on + button:

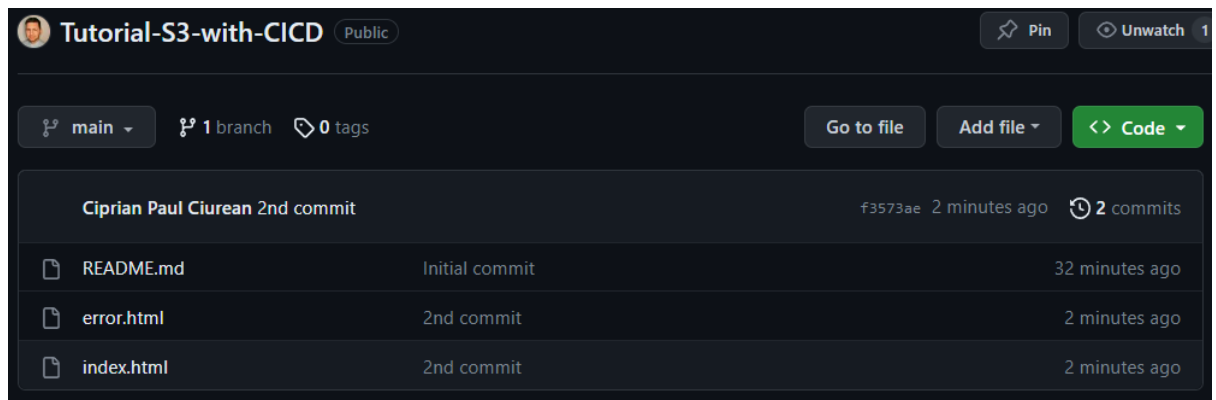


3. Type a message and press **Commit**.



4. Press **Sync Changes** and **OK** to commit the changes.

Now if you go online into your github account, you will find the 2 new files added to your repository.



Well done! We created our repository online where we will store our website files and this is where we will connect our pipeline to make the changes to our website easier.

Step 5: CI/CD Pipeline with AWS CodePipeline

5.1 Create the Pipeline:

Choose pipeline settings Info

Pipeline settings

Pipeline name
Enter the pipeline name. You cannot edit the pipeline name after it is created.

No more than 100 characters

Service role

☒ **New service role**
Create a service role in your account

☐ **Existing service role**
Choose an existing service role from your account

Role name

Type your service role name

☒ Allow AWS CodePipeline to create a service role so it can be used with this new pipeline

► **Advanced settings**

Cancel Next

5.2 Add source stage:

Source

Source provider
This is where you stored your input artifacts for your pipeline. Choose the provider and then provide the connection details.

New GitHub version 2 (app-based) action
To add a GitHub version 2 action in CodePipeline, you create a connection, which uses GitHub Apps to access your repository. Use the options below to choose an existing connection or create a new one. [Learn more](#)

Connection
Choose an existing connection that you have already configured, or create a new one and then return to this task.

or Connect to GitHub

5.3 Create a connection:

Create a connection [Info](#)

Create GitHub App connection [Info](#)

Connection name

► **Tags - optional**

[Connect to GitHub](#)


5.4 Connect to GitHub:

Connect to GitHub

GitHub connection settings [Info](#)

Connection name

GitHub Apps
GitHub Apps create a link for your connection with GitHub. To start, install a new app and save this connection.



 or [Install a new app](#)

► **Tags - optional**

[Connect](#)

Install a new app:

Permissions

- ✓ Read access to issues and metadata
- ✓ Read and write access to administration, code, commit statuses, and pull requests


Repository access

☒ **All repositories**
This applies to all current *and* future repositories owned by the resource owner.
Also includes public repositories (read-only).

☐ **Only select repositories**
Select at least one repository.
Also includes public repositories (read-only).

Selected 1 repository:

×

**Ready to connect**
Your GitHub connection is ready for use.

Repository name

Choose a repository in your GitHub account.

You can type or paste the group path to any project that the provided credentials can access. Use the format 'group/subgroup/project'.

Branch name

Choose a branch of the repository.

Change detection options

☒ **Start the pipeline on source code change**
Automatically starts your pipeline when a change occurs in the source code. If turned off, your pipeline only runs if you start it manually or on a schedule.

Output artifact format

Choose the output artifact format.

☒ **CodePipeline default**
AWS CodePipeline uses the default zip format for artifacts in the pipeline. Does not include Git metadata about the repository.

☐ **Full clone**
AWS CodePipeline passes metadata about the repository that allows subsequent actions to do a full Git clone. Only supported for AWS CodeBuild actions.

We can skip the building page as we build our files in VSC.

5.5 Add deploy stage:

Deploy

Deploy provider
Choose how you deploy to instances. Choose the provider, and then provide the configuration details for that provider.

Amazon S3 ▼

Region
US East (N. Virginia) ▼

Bucket
paulciurean.com ✕

Deployment path - *optional*

☒ Extract file before deploy
The deployed artifact will be unzipped before deployment.

► Additional configuration

Cancel Previous Next

5.6 Review the pipeline:

Step 4: Add deploy stage

Deploy action provider

Deploy action provider
Amazon S3
Extract
true
BucketName
demo1910.com

Cancel Previous Create pipeline

AWS test the pipeline and we receive a **SUCCESS**

The screenshot shows the AWS CodePipeline console for a pipeline named 'Tutorial-S3-website'. At the top, there are buttons for 'Notify', 'Edit', 'Stop execution', 'Clone pipeline', and 'Release change'. The pipeline execution ID is '2e2850fe-abb0-4b1e-b84c-76f74e8b1b13'. The pipeline consists of two stages: 'Source' and 'Deploy'. The 'Source' stage is marked as 'Succeeded' and used 'GitHub (Version 2)' provider. The 'Deploy' stage is also marked as 'Succeeded' and used 'Amazon S3' provider. A 'Disable transition' button is visible between the stages. On the right side, there are two green checkmarks indicating success.

CodePipeline is now pulling any changes from GitHub and deploying them to the bucket that the website is hosted in. Looking at the S3 bucket you can see all the files have been updated at the same time. This is because CodePipeline deploys the entire repository rather than just individual changed files.

The screenshot shows the Amazon S3 console 'Objects' view for a bucket. It displays a list of three objects: 'error.html', 'index.html', and 'README.md'. All objects were last modified on October 20, 2023, at 15:57:12 (UTC+01:00). The sizes are 337.0 B, 418.0 B, and 127.0 B respectively, all using the 'Standard' storage class. The table has columns for Name, Type, Last modified, Size, and Storage class.

<input type="checkbox"/>	Name	Type	Last modified	Size	Storage class
<input type="checkbox"/>	error.html	html	October 20, 2023, 15:57:12 (UTC+01:00)	337.0 B	Standard
<input type="checkbox"/>	index.html	html	October 20, 2023, 15:57:12 (UTC+01:00)	418.0 B	Standard
<input type="checkbox"/>	README.md	md	October 20, 2023, 15:57:12 (UTC+01:00)	127.0 B	Standard

Step 6: Create the Lambda Function

6.1 Create Function:

Create function [Info](#)

AWS Serverless Application Repository applications have moved to [Create application](#).

☒ Author from scratch
Start with a simple Hello World example.

☐ Use a blueprint
Build a Lambda application from sample code and configuration presets for common use cases.

☐ Container image
Select a container image to deploy for your function.

Basic information

Function name
Enter a name that describes the purpose of your function.

Use only letters, numbers, hyphens, or underscores with no spaces.

Runtime [Info](#)
Choose the language to use to write your function. Note that the console code editor supports only Node.js, Python, and Ruby.

Python 3.9

Architecture [Info](#)
Choose the instruction set architecture you want for your function code.

☒ x86_64
☐ arm64

Permissions [Info](#)
By default, Lambda will create an execution role with permissions to upload logs to Amazon CloudWatch Logs. You can customize this default role later when adding triggers.

[Change default execution role](#)

6.2 Copy the python code in lambda function

lambda_function × Environment Var × +

```
7         if record['s3']['object']['key'] == "index.html":
8             invalidation = cloudfront_client.create_invalidation(
9                 DistributionId = "E2EN7DX4ZHL16K",
10                 InvalidationBatch = {
11                     'Paths': {
12                         "Quantity": 1,
13                         "Items": ["/"]
14                     },
15                     "CallerReference": str(record['eventTime'])
16                 }
17             )
18             return {
19                 'statusCode': 200,
20                 'body': json.dumps("Cloudfront invalidation successfully created")
21             }
22         return {
23             'statusCode': 204,
24             'body': json.dumps("No file called index.html therefore no invalidation created")
25         }
```

PS: You can find the code in my GitHub but remember to insert your CloudFront distribution on line 5.

6.2 Add permissions to Lambda to be able to create invalidation


6.2.1 In Lambda go to Configuration -> Permission

Execution role Refresh Edit View role document

Role name
Invalidation-Function-role-o94hgwo1 [Link](#)

Resource summary

To view the resources and actions that your function has permission to access, choose a service.

 Amazon CloudWatch Logs
3 actions, 2 resources

6.2.2 Add permissions

<input type="checkbox"/>	Policy name Link	Type	Attached entities
<input type="checkbox"/>	AWSLambdaBasicExecutionRole-b871d338-0a56-4ab3-913...	Customer managed	1
<input type="checkbox"/>	Invalidation-policy	Customer inline	0

Invalidation-policy Copy JSON Edit

```
1 {  
2   "Version": "2012-10-17",  
3   "Statement": [  
4     {  
5       "Sid": "AllowDistributionInvalidation",  
6       "Effect": "Allow",  
7       "Action": "cloudfront:CreateInvalidation",  
8       "Resource": "arn:aws:cloudfront::753656081091:distribution/E26OLYTU4CRYT1"  
9     }  
10  ]  
11 }
```

Create a new inline policy. You can find the JSON for it in [GitHub](#)

6.3 Create a test event

Test event [info](#) Save Test

To invoke your function without saving an event, configure the JSON event, then choose Test.

Test event action
☒ Create new event ☐ Edit saved event

Event name

Maximum of 25 characters consisting of letters, numbers, dots, hyphens and underscores.

Event sharing settings
☒ Private
This event is only available in the Lambda console and to the event creator. You can configure a total of 10. [Learn more](#)
☐ Shareable
This event is available to IAM users within the same account who have permissions to access and use shareable events. [Learn more](#)

Template - optional

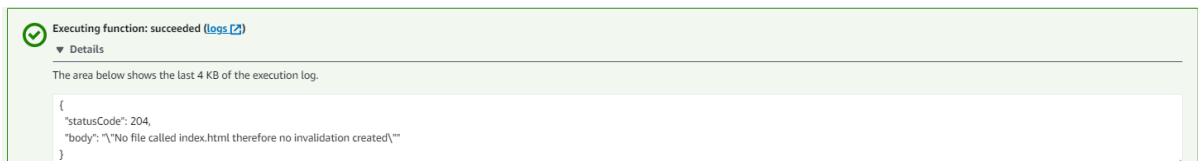
Make sure you change the following:

Line 10 - replace “example” with your account number

Line 25 - replace “example” with your account number

Line 27 – replace “example-bucket” with your bucket name


6.3 Test the function



6.4 Trigger lambda with S3 events

Add trigger

Trigger configuration [Info](#)

 **S3**
aws asynchronous storage

Bucket
Please select the S3 bucket that serves as the event source. The bucket must be in the same region as the function.
 ✕ ↺
Unable to verify bucket region: Access Denied

Event types
Select the events that you want to have trigger the Lambda function. You can optionally set up a prefix or suffix for an event. However, for each bucket, individual events cannot have multiple configurations with overlapping prefixes or suffixes that could match the same object key.


PUT ✕

To make sure the invalidation is triggered only when you change the .html file we add the following to SUFFIX


Suffix - optional

Enter a single optional suffix to limit the notifications to objects with keys that end with matching characters.

Recursive invocation

If your function writes objects to an S3 bucket, ensure that you are using different S3 buckets for input and output. Writing to the same bucket increases the risk of creating a recursive invocation, which can result in increased Lambda usage and increased costs. [Learn more](#) 

- ☒ I acknowledge that using the same S3 bucket for both input and output is not recommended and that this configuration can cause recursive invocations, increased Lambda usage, and increased costs.

Lambda will add the necessary permissions for AWS S3 to invoke your Lambda function from this trigger. [Learn more](#)  about the Lambda permissions model.

Cancel

Add

6.5 Final test

I changed the code in index.html file and the lambda created the invalidation successfully:

○ [JF0WUUSNIX6O70YVY2SB1XMICC](#) In progress October 21, 2023 at 10:11:04 AM UTC

The website was updated within seconds:

Congratulations, you've launched a static website on Amazon S3

This website was securely accessed through Amazon CloudFront

Final test 21.10.2023

7. Conclusion

In this tutorial, we've seen how you can keep your data safe. S3 buckets, which store your website files, should never be left open to the public (unless we are required to do so) because there are always people looking for exposed data. We can protect them using methods like Origin Access Control and by making sure services don't need to go over the public internet to access these files.

We've also learned about a handy way to update your website automatically. This is done using AWS CodePipeline, but there are other services like AWS CodeCommit (in our case GitHub) and AWS CodeDeploy that help you manage your code and deploy it to different parts of AWS.

Lastly, we explored an easy way to automate tasks by using AWS Lambda. This keeps your website up to date without having to manually make changes to CloudFront.

So, in simple terms, our setup emphasizes the need for safety, easy updates, and automation in your AWS projects. These practices will make your system more secure, efficient, and reliable.

Thank you for following this tutorial. In case you need any help or guidance don't hesitate to contact me.
www.paulciurean.com